

Exercises3

November 13, 2020

1 Exercise set 3

1.1 Answered to all questions (1-5)

1.2 Question 1

$$S_W = C_0 + C_1 = \begin{pmatrix} 3 & 0 \\ 0 & 1 \end{pmatrix} + \begin{pmatrix} 3 & -1 \\ -1 & 1 \end{pmatrix} = \begin{pmatrix} 6 & -1 \\ -1 & 2 \end{pmatrix}$$

$$S_W^{-1} = \frac{1}{6 * 2 - 1} \begin{pmatrix} 6 & -1 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} \frac{2}{11} & \frac{1}{11} \\ \frac{1}{11} & \frac{6}{11} \end{pmatrix}$$

Since the scale of projection vector \mathbf{w} doesn't matter, we can write

$$\mathbf{w} = S_W^{-1}(\mu_1 - \mu_0) = \begin{pmatrix} 6 & -1 \\ -1 & 2 \end{pmatrix} \left(\begin{pmatrix} 1 \\ 2 \end{pmatrix} - \begin{pmatrix} 1 \\ 1 \end{pmatrix} \right) = \begin{pmatrix} 1 \\ 6 \end{pmatrix}$$

1.3 Question 2

When classification problem is defined as a likelihood ratio test, I got following thresholds for x (Derivation of thresholds is omitted, since it is long and messy): Classify projected sample x in class 1 if

$$c - \sqrt{\frac{2\log(\sigma_1\sigma_2^{-1}) - \sigma_2^{-2}\mu_2^2 - \sigma_1^{-2}\mu_1^2 + \frac{(\sigma_2^{-2}\mu_2 - \sigma_1^{-2}\mu_1)^2}{\sigma_2^{-2} - \sigma_1^{-2}}}{\sigma_2^{-2} - \sigma_1^{-2}}} < x < c + \sqrt{\frac{2\log(\sigma_1\sigma_2^{-1}) - \sigma_2^{-2}\mu_2^2 - \sigma_1^{-2}\mu_1^2 + \frac{(\sigma_2^{-2}\mu_2 - \sigma_1^{-2}\mu_1)^2}{\sigma_2^{-2} - \sigma_1^{-2}}}{\sigma_2^{-2} - \sigma_1^{-2}}}$$

where

$$c = \frac{\sigma_2^{-2}\mu_2 - \sigma_1^{-2}\mu_1}{\sigma_2^{-2} - \sigma_1^{-2}}$$

$$\mu_1 = \mathbf{w}^T \boldsymbol{\mu}_0$$

$$\sigma_1^2 = \mathbf{w}^T \mathbf{C}_0 \mathbf{w}$$

$$\mu_2 = \mathbf{w}^T \boldsymbol{\mu}_1$$

$$\sigma_2^2 = \mathbf{w}^T \mathbf{C}_1 \mathbf{w}$$

In the context of Question 1, numeric values of these parameters are:

$$\mu_1 = \mathbf{w}^T \boldsymbol{\mu}_0 = 7$$

$$\sigma_1^2 = \mathbf{w}^T \mathbf{C}_0 \mathbf{w} = 39$$

$$\mu_2 = \mathbf{w}^T \boldsymbol{\mu}_1 = 13$$

$$\sigma_2^2 = \mathbf{w}^T \mathbf{C}_1 \mathbf{w} = 27$$

By applying these to the formula given above, obtained thresholds are:

$$9.30937 < x < 43.6906$$

By using the \mathbf{w} derived in Q1, projected x is:

$$x = (1 \ 6) \begin{pmatrix} 1 \\ 2 \end{pmatrix} = 13$$

Therefore, the point \mathbf{x} is classified to class 1.

1.4 Question 3

```
[3]: # Imports
import time
import numpy as np
from skimage.io import imread_collection
from skimage.transform import rescale, resize, downscale_local_mean
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import cv2
from sklearn.metrics import accuracy_score

from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import MinMaxScaler

root = '/home/tuomas/Python/DATA.ML.200/Ex3/'
images = []
labels = []
for i in range(0,9):
    fn = '0000{}/*.jpg'.format(i)
    print(fn)
    imgs = imread_collection(root + fn)
    images.append(np.array(imgs, dtype='object'))
    labels.append( np.ones(len(imgs)) * i )

#images = np.array(images, dtype='object')
images = np.concatenate(images)
labels = np.concatenate(labels).astype('uint8')
```

```
00000/*.jpg
00001/*.jpg
00002/*.jpg
00003/*.jpg
00004/*.jpg
00005/*.jpg
00006/*.jpg
00007/*.jpg
00008/*.jpg
```

```
[4]: # Preprocess the images
     imgs_processed = []
     scaler = MinMaxScaler()
     for img in images:
         # Resize & vectorize the image
         img = cv2.resize(img, (32,32)).ravel()
         # Scale sample to (0,1)
         # Since MinMaxScaler scales data featurewise, I transposed the row vector
         → into column vector
         # so the samplewise scaling is performed
         img_T = img[... ,None]
         scaler.fit(img_T)
         img_T = scaler.transform(img_T)

         imgs_processed.append(img_T.ravel())

     imgs_processed = np.array(imgs_processed).astype('float32')
```

```
[5]: # Create training and testing sets
     trainX, testX, trainY, testY = train_test_split(imgs_processed,
                                                    labels,
                                                    test_size=0.3)
```

```
[4]: # Train, test & evaluate given models
     models = [KNeighborsClassifier(n_neighbors=3),
               LinearDiscriminantAnalysis(solver='svd'),
               LogisticRegression(max_iter=10000),
               SVC(kernel='linear'),
               SVC(kernel='rbf'),
               RandomForestClassifier(n_estimators=20)
               ]

     accuracy = []
     tr_time = []
     tst_time = []
     for model in models:
         # Training
```

```

print('Training {} ...'.format(model.__class__.__name__))
start = time.time()
model.fit(trainX, trainY)
tr_time.append( time.time() - start )
# Testing
print('Testing {} ...'.format(model.__class__.__name__))
start = time.time()
predY = model.predict(testX)
tst_time.append( time.time() - start )
# Evaluating
print('Evaluating {} ...'.format(model.__class__.__name__))
accuracy.append( accuracy_score(testY, predY) )

```

```

Training KNeighborsClassifier ...
Testing KNeighborsClassifier ...
Evaluating KNeighborsClassifier ...
Training LinearDiscriminantAnalysis ...
Testing LinearDiscriminantAnalysis ...
Evaluating LinearDiscriminantAnalysis ...
Training LogisticRegression ...
Testing LogisticRegression ...
Evaluating LogisticRegression ...
Training SVC ...
Testing SVC ...
Evaluating SVC ...
Training SVC ...
Testing SVC ...
Evaluating SVC ...
Training RandomForestClassifier ...
Testing RandomForestClassifier ...
Evaluating RandomForestClassifier ...

```

```

[5]: # Print the statistics
model_names = ['3-NN', 'LDA', 'LogReg', 'SVM linear', 'SVM rbf', 'Random forest']
print('Training set size (batch) = {}'.format(trainX.shape[0]))
print('Test set size = {}\n'.format(testX.shape[0]))
for i in range(6):
    print('Results for {}'.format(model_names[i]))
    print('-----')
    print('Accuracy is {} %'.format(round(accuracy[i]*100, 2)))
    print('Training time / batch {} s'.format(round(tr_time[i], 2)))
    print('Test time / sample {} ms\n'.format(round(tst_time[i]*1000/testX.
↪shape[0], 3)))

```

```

Training set size (batch) = 6237
Test set size = 2673

```

Results for 3-NN

```
-----  
Accuracy is 88.78 %  
Training time / batch 1.65 s  
Test time / sample 18.116 ms
```

Results for LDA

```
-----  
Accuracy is 80.96 %  
Training time / batch 11.56 s  
Test time / sample 0.006 ms
```

Results for LogReg

```
-----  
Accuracy is 94.28 %  
Training time / batch 32.31 s  
Test time / sample 0.006 ms
```

Results for SVM linear

```
-----  
Accuracy is 95.06 %  
Training time / batch 32.63 s  
Test time / sample 5.539 ms
```

Results for SVM rbf

```
-----  
Accuracy is 91.99 %  
Training time / batch 70.33 s  
Test time / sample 11.673 ms
```

Results for Random forest

```
-----  
Accuracy is 91.81 %  
Training time / batch 3.52 s  
Test time / sample 0.007 ms
```

1.5 Question 4

By previous question, it seems that SVM linear had the best performance

```
[6]: # Concatenate train and test set from previous exercise  
from sklearn.model_selection import cross_val_score  
trainX2 = np.concatenate((trainX, testX))  
trainY2 = np.concatenate((trainY, testY))  
model = SVC(kernel='linear')  
cv_scores = cross_val_score(model, trainX2, trainY2, cv=5, verbose=2)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
```

```

[CV] ...
[CV] ... , total= 51.5s
[CV] ...

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 51.5s remaining: 0.0s

[CV] ... , total= 51.0s
[CV] ...
[CV] ... , total= 51.1s
[CV] ...
[CV] ... , total= 51.1s
[CV] ...
[CV] ... , total= 51.0s

[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 4.3min finished

```

```

[8]: print('Mean of accuracies = {}'.format(np.mean(cv_scores)))
     print('Standard deviation of accuracies = {}'.format(np.std(cv_scores)))

```

```

Mean of accuracies = 0.9584736251402919
Standard deviation of accuracies = 0.0015470312853075455

```

1.6 Question 5

```

[11]: # Resize all images to 32 x 32
      images_resized = []
      for img in images:
          images_resized.append(cv2.resize(img, (32,32)))

      images_resized = np.array(images_resized)

```

```

[12]: # Create training and testing sets.
      from tensorflow.keras.utils import to_categorical
      trainX, testX, trainY, testY = train_test_split(images_resized,
                                                         labels,
                                                         test_size=0.15)

      # One-hot encoding
      trainY = to_categorical(trainY, num_classes=9)
      testY_cat = to_categorical(testY, num_classes=9)

```

```

[13]: # Build the model
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, BatchNormalization, Dropout, Flatten, Dense

      layers=[Conv2D(filters=32, kernel_size = 3, activation='relu', input_shape = (32,32,3)),
              BatchNormalization(),

```

```

        Conv2D(filters=32, kernel_size = 3, activation='relu'),
        BatchNormalization(),
        Conv2D(32, kernel_size = 5, strides=2, padding='same',
→activation='relu'),
        BatchNormalization(),
        Dropout(0.4),
        Conv2D(64, kernel_size = 3, activation='relu'),
        BatchNormalization(),
        Conv2D(64, kernel_size = 3, activation='relu'),
        BatchNormalization(),
        Conv2D(64, kernel_size = 5, strides=2, padding='same',
→activation='relu'),
        BatchNormalization(),
        Dropout(0.4),
        Conv2D(128, kernel_size = 4, activation='relu'),
        BatchNormalization(),
        Flatten(),
        Dropout(0.4),
        Dense(9, activation='softmax'])

cnn_model = Sequential(layers)
cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',
→metrics=['accuracy'])
cnn_model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
batch_normalization (Batch Normalization)	(None, 30, 30, 32)	128
conv2d_1 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 28, 28, 32)	128
conv2d_2 (Conv2D)	(None, 14, 14, 32)	25632
batch_normalization_2 (Batch Normalization)	(None, 14, 14, 32)	128
dropout (Dropout)	(None, 14, 14, 32)	0
conv2d_3 (Conv2D)	(None, 12, 12, 64)	18496
batch_normalization_3 (Batch Normalization)	(None, 12, 12, 64)	256

conv2d_4 (Conv2D)	(None, 10, 10, 64)	36928

batch_normalization_4 (Batch Normalization)	(None, 10, 10, 64)	256

conv2d_5 (Conv2D)	(None, 5, 5, 64)	102464

batch_normalization_5 (Batch Normalization)	(None, 5, 5, 64)	256

dropout_1 (Dropout)	(None, 5, 5, 64)	0

conv2d_6 (Conv2D)	(None, 2, 2, 128)	131200

batch_normalization_6 (Batch Normalization)	(None, 2, 2, 128)	512

flatten (Flatten)	(None, 512)	0

dropout_2 (Dropout)	(None, 512)	0

dense (Dense)	(None, 9)	4617
=====		
Total params: 331,145		
Trainable params: 330,313		
Non-trainable params: 832		

```
[14]: # Train the model
      cnn_model.fit(trainX, trainY, epochs=10, batch_size=64, verbose=1)
```

```
Epoch 1/10
119/119 [=====] - 6s 48ms/step - loss: 1.4053 -
accuracy: 0.5504
Epoch 2/10
119/119 [=====] - 1s 10ms/step - loss: 0.3055 -
accuracy: 0.8946
Epoch 3/10
119/119 [=====] - 1s 10ms/step - loss: 0.1569 -
accuracy: 0.9477
Epoch 4/10
119/119 [=====] - 1s 10ms/step - loss: 0.0982 -
accuracy: 0.9663
Epoch 5/10
119/119 [=====] - 1s 10ms/step - loss: 0.0723 -
accuracy: 0.9781
Epoch 6/10
119/119 [=====] - 1s 10ms/step - loss: 0.0595 -
accuracy: 0.9795
Epoch 7/10
119/119 [=====] - 1s 10ms/step - loss: 0.0536 -
```



```
accuracy: 0.9814
Epoch 8/10
119/119 [=====] - 1s 10ms/step - loss: 0.0438 -
accuracy: 0.9852
Epoch 9/10
119/119 [=====] - 1s 10ms/step - loss: 0.0302 -
accuracy: 0.9901
Epoch 10/10
119/119 [=====] - 1s 10ms/step - loss: 0.0314 -
accuracy: 0.9880
```

```
[14]: <tensorflow.python.keras.callbacks.History at 0x7fea1484cf10>
```

```
[15]: # Evaluate the model
test_loss, test_acc = cnn_model.evaluate(testX, testY_cat, verbose=2)
print("Accuracy of CNN = {}".format(test_acc))
```

```
42/42 - 2s - loss: 0.0337 - accuracy: 0.9933
Accuracy of CNN = 0.9932684898376465
```