

DATA.ML.200 Pattern Recognition and Machine Learning

Exercise Set 4: November 16–November 20, 2020

- Exercises consist of both **pen&paper** and **python** assignments.
- Prepare a single PDF and return to Moodle on Friday, November 20th at 23:55 at the latest.
- Mark on 1st page which exercises you did.

1. **pen&paper** Compute the gradient of the log-loss.

In the lectures we defined the *logistic loss function*:

$$\ell(\mathbf{w}) = \sum_{n=0}^{N-1} \ln(1 + \exp(-y_n \mathbf{w}^T \mathbf{x}_n)), \quad (1)$$

and computed its gradient $\frac{\partial \ell(\mathbf{w})}{\partial \mathbf{w}}$. Here, $\mathbf{x}_n \in \mathbf{R}^P$ and $y_n \in \{-1, 1\}$ are the inputs and labels for the samples $n = 0, 1, \dots, N - 1$, and $\mathbf{w} \in \mathbf{R}^P$ are the model parameters to be learnt.

The L_2 -regularized logistic loss is defined by:

$$\ell(\mathbf{w}) = \sum_{n=0}^{N-1} \ln(1 + \exp(-y_n \mathbf{w}^T \mathbf{x}_n)) + C \cdot \mathbf{w}^T \mathbf{w}, \quad (2)$$

with $C \geq 0$ the regularization strength parameter. Compute the gradient of the regularized loss.

Hint: For finding the gradients of vector functions, check the document at http://www.kamperh.com/notes/kamper_matrixcalculus13.pdf

2. **pen&paper** Consider the following Keras code defining a convolutional neural network.

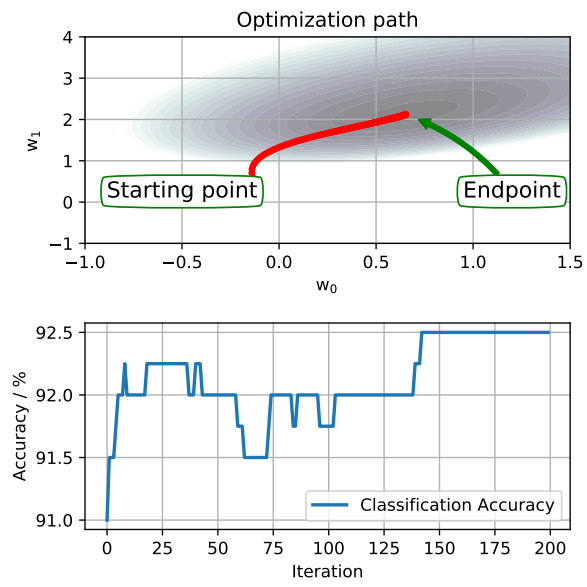
```
N = 10          # Number of feature maps
w, h = 5, 5     # Conv. window size

model.add(Conv2D(N, (w, h),
                  input_shape=(64, 64, 3),
                  activation = 'relu',
                  padding = 'same'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(N, (w, h),
                  activation = 'relu',
                  padding = 'same'))
model.add(MaxPooling2D((2, 2)))

model.add(Flatten())
model.add(Dense(2, activation = 'sigmoid'))
```

- a) Draw a diagram of the network similar to the one at the bottom of slide 14 in <http://www.cs.tut.fi/courses/SGN-41007/slides/Lecture6.pdf>
- b) Compute the number of parameters of the network at each layer (and explain why).
3. **python** Implement gradient descent for log-loss.
- a) Implement a log-loss minimization algorithm for the loss of Equation (1). You may use the template provided by the teaching assistant.
- b) Apply the code for the data downloaded from
https://github.com/mahehu/SGN-41007/tree/master/exercises/Ex5/log_loss_data.zip
The data is in CSV format. Load X and y using `numpy.loadtxt`.
- c) Plot the path of w over 100 iterations and check the accuracy (see plots below).



4. `python` Define the network in Keras.

There are two options for this task. Either a) or b) will be enough.

- a) Study from the slides and Tensorflow/Keras documentation how to instantiate a pretrained **Mobilenet V2** model with input shape $64 \times 64 \times 3$ and 9 outputs. Last layer should be a softmax layer.
- b) Implement a neural network such that `tf.keras.model.summary()` gives the following output. Last layer should be a softmax layer.

<code>model.summary()</code>		
<code>Layer (type)</code>	<code>Output Shape</code>	<code>Param #</code>
=====		
<code>conv2d_49 (Conv2D)</code>	<code>(None, 64, 64, 32)</code>	<code>2432</code>
<code>max_pooling2d_47 (MaxPooling)</code>	<code>(None, 16, 16, 32)</code>	<code>0</code>
<code>conv2d_50 (Conv2D)</code>	<code>(None, 16, 16, 32)</code>	<code>25632</code>
<code>max_pooling2d_48 (MaxPooling)</code>	<code>(None, 4, 4, 32)</code>	<code>0</code>
<code>flatten_15 (Flatten)</code>	<code>(None, 512)</code>	<code>0</code>
<code>dense_29 (Dense)</code>	<code>(None, 100)</code>	<code>51300</code>
<code>dense_30 (Dense)</code>	<code>(None, 9)</code>	<code>909</code>
=====		
<code>Total params: 80,273</code>		
<code>Trainable params: 80,273</code>		
<code>Non-trainable params: 0</code>		

5. `python` *Compile and train the net.*

- a) Compile the network of Question 4 above.
- b) Train the model with the GTSRB dataset from last week.

Use the following parameters:

- **Loss:** categorical crossentropy (same thing as log loss; see previous exercises)
- **Optimizer:** Adam or Stochastic gradient descent
- **Minibatch size:** 32
- **Number of epochs:** 20

Also add the parameter `metrics=['accuracy']` as an argument of `model.compile` and give the test data to training algorithm `model.fit(..., validation_data = [X_test, y_test])` Then, the optimizer will report the test error every epoch.